

Introduction to Cluster Scheduler Nomad

Easily Deploy Applications at Any Scale

2018

SHLUG

- Su Yan
- Pure FP, Type Theory
- Infrastructure

Goal

- Cluster Scheduler Concepts
- Nomad as an alternative to k8s/mesos

Overview & History

What's Cluster Scheduler

- a tool for managing a cluster of machines and running applications on them

Why use it

- Hard to change things in production
- Snowflakes in production
 - run 30 instances of app Foo
 - run only 1 instance of app Bar (problematic otherwise)
 - run X & Y together
 - ad-hoc scripts (rendering configs, monitoring)
- Resource Utilization - multiple apps per server

- Monoliths - high application complexity
- Microservice - high operational complexity

- Google Borg (Omega)
- Tupperware (Facebook)
- Apache Mesos (2010) (UC Berkely, Twitter)
- Apache Aurora (2010 , OSS 2013) (Twitter)
- Mesosphere, YARN, ECS, Rancher, Nomad, Kubernetes...

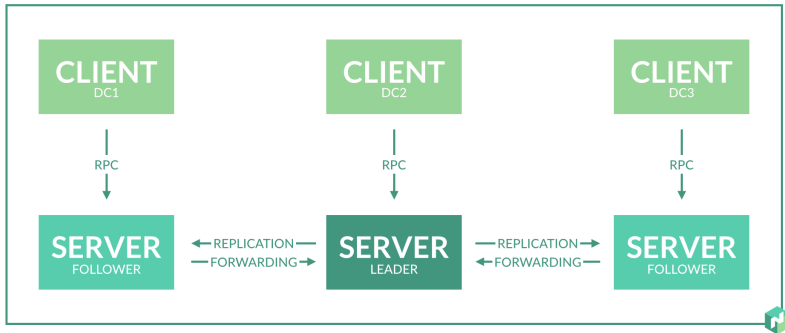
Nomad

- Declarative jobs
- Infrastructure as Code
- Consul & Vault Integration
- ACL
- Operationally Simple & Scalable
 - Single Binary
 - No Dependencies
 - Highly Available
 - Multi-DC/Region Support

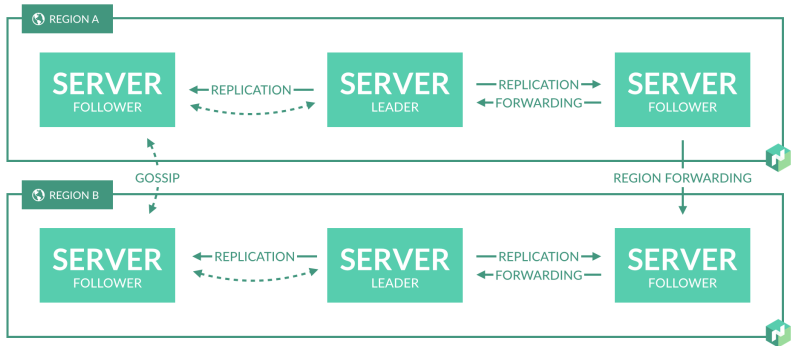
- Rolling updates, rollbacks
- Blue/Green depoyments
- Automatic failure handling
- Fexible workloads
 - linux, windows, mac
 - container (docker, rkt, lxc), VM, raw process
 - service, batch (spark)

Architecture

Single Region



Global



Job specification

```
job
  \_ group[]
      \_ task[]
```

- HCL or JSON
- Job – > Task Group[] – > Task[]
- Job: datacenters, region, type(service/batch/system), update/deployment strategy, priority, constraints, meta, ...
- Task Group: count, meta, ...
- Task: driver, config, resources, logs, meta, ...
- Resources: cpu, disk, iops, memory, network (ports, mbits)

```
job "docs" {  
  region = "us"  
  datacenters = ["us-west-1", "us-east-1"]  
  type = "service"  
  update {  
    stagger      = "30s"  
    max_parallel = 2  
  }  
  group "webs" {  
    count = 5  
    task "frontend" {...}  
  }  
}
```

```
task "frontend" {  
    driver = "docker"  
    config {  
        image = "hashicorp/web-frontend"  
    }  
    ...  
}
```

Job Spec - Service Registration (Consul)

```
task "frontend" {  
  ...  
  service {  
    port = "http"  
    check {  
      type      = "http"  
      path      = "/health"  
      interval  = "10s"  
      timeout   = "2s"  
    }  
  }  
}
```

```
task "frontend" {  
    ...  
    env {  
        "DB_HOST" = "db01.example.com"  
        "DB_USER" = "web"  
        "DB_PASS" = "loremipsum"  
    }  
    ...  
}
```

```
task "frontend" {  
    ...  
    resources {  
        network {  
            mbits = 100  
            port "http" {}  
            port "https" {  
                static = 443  
            }  
        }  
    }  
    ...  
}
```

- `NOMAD_META_{key} = {value}`
- `NOMAD_CPU_LIMIT`
- `NOMAD_MEMORY_LIMIT`
- `NOMAD_IP`
- `NOMAD_PORT_{label}`

- artifact
- check_restart
- constraint
- dispatch_payload
- env
- ephemeral_disk
- group
- job
- logs
- meta
- network
- parameterized
- periodic
- resources
- restart
- service
- task
- template
- update
- vault

- TLS
- Vault integration
- ACL
- Sentinel (policy as code) (paid)

- /v1/acl
- /v1/jobs
- /v1/nodes
- /v1/allocations
- /v1/evaluations
- /v1/deployments
- /v1/agent/...
- /v1/status

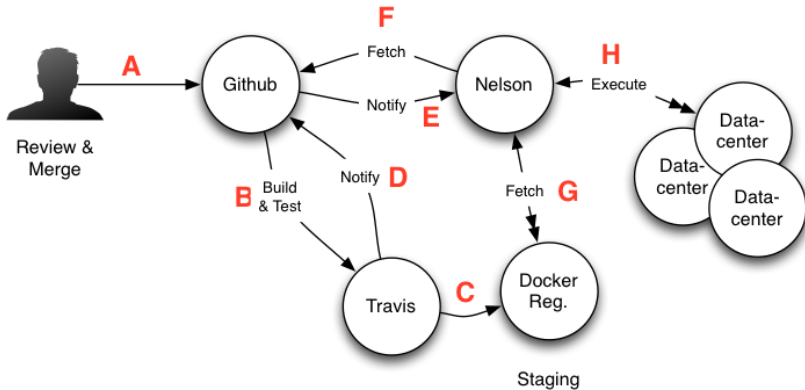
- Stateful Application*
- Volume Plugin (CSI)
- Network Plugin (CNI)
- Priorities (preemption)
- Quotas

built-in (last year)

demo: <https://demo.nomadproject.io/>

- Serverless (OpenFaaS plugin)
- Spark (fork), Apache Heron
- Nelson (workflow system) (getnelson.github.io/nelson)
 - Fully integrated with GitHub or GitHub Enterprise.
 - Developer-driven, automated build & release workflow revisioned as code.
 - Can deploy applications to any number of datacenters.
 - State of the art runtime routing via Envoy.
 - Integrated support for alert definition and propagation via Prometheus.
 - Utilizes secure introduction for safe distribution of credentials from Vault.

Nelson



Start small/simple

Servers

- 3 node
- run nomad, consul & vault server

Workers

- provision nomad and consul agent
- consul agent need 1 consul server address to bootstrap
- nomad connect to servers via consul

Consul DNS

- do use for trivial use cases
- more sophisticated solutions
 - istio/envoy/linkerd
 - fat client libraries

Server upgrade

- rolling upgrade
- make sure server stand up before starting next one

Use ACL, TLS, Vault from the very beginning

- Run bigger and less servers, easier to do experimentation
- Code review!
- Survive mandatory security update in public cloud
- Portable (k8s, mesos, etc.)

Thanks

Questions?